

Experimentation with Clustering Algorithms

Corinne Curcie and Shivam Naik

March 5, 2017

Abstract

Clustering algorithms are a popular tool for making sense of big data. Our project involved implementing various algorithms, focusing specifically on high-dimensional data, to gain better understanding .

1. Introduction
2. Related Work
3. Algorithms Implemented

K-Means

One of the algorithms we implemented is a variant on K-Means. Before getting into the details of that algorithm, we will go over the basics of K-Means, since we also used it as a benchmark when scoring preformances of algorithms we implemented. The well-known K-Means algorithm starts by choosing k points in the dataset to be the initial cluster center points, and then updates on an “expectation-maximization mechanism” (EM).

The “E-step” is the cluster assignment step, where points are labelled with a cluster based on which of the center points they are closest to. The “M-step” is the update step, where the k cluster center points are recalculated to be some average of all of the points that were labeled as belonging to that cluster during the previous “E-step” round. The algorithm stops when the change in centers from one round to the next is less than some predetermined threshold. The objective function that K-Means tries to minimize is the sum of squared Euclidean distances from each point to its assigned cluster center:

$$\arg \min_{\mathbf{S}} \sum_{i=1}^k \sum_{\mathbf{x} \in S_i} \|\mathbf{x} - \boldsymbol{\mu}_i\|^2$$

K-Subspaces

We were intrigued by a 2009 paper that proposed a “K-Subspaces” algorithm similar to K-Means (Wang, Ding, and Li 2009). Instead of using Euclidean distances for our objective function, multiple distance measures are used during the “E-step” to determine which centers the points are closest to. The authors of the algorithm decided to focus on three possible subspaces – 1D lines, 2D planes, and 3D spheres – and determines distance functions based on those subspaces. By calculating all three distances for each pair of point and cluster center, the algorithm is better able to determine when a point is within a cluster of a non-standard space. To extend this performance not only to the entire dataset but also subspaces within the dataset, PCA is used during distance calculations to isolate meaningful subspaces

K-Subspaces Initialization

For initializing cluster centers before beginning the EM steps, we used the standard “K-means++” algorithm, which probabilistically selects initial clusters. Specifically, the algorithm goes through the following steps: (1) Select an input data point uniformly at random to be the first center. (2) Calculate the distance between every point x and its nearest center point that has already been chosen. (3) Select a new data point as a next center, using a probability distribution where a point is chosen with probability proportional to the distance calculated in step 2. (4) Repeat the steps 2 and 3 until k centers have been chosen. The “K-means++” initialization is also used in scikit-learns implementation of K-Means.

Cluster Assignment Step

In our cluster assignment step is when we use the new distance measures proposed by the authors. Instead of finding the direct distance to the center of a cluster, they aim to estimate a *surface* of a cluster by using perpendicular distances. For the following equations, let x be a point and C_k represent a cluster where c_k is the center point of the cluster.

Line Distance:

$$Dist(\mathbf{x}, C_k) = \|\mathbf{x}^\perp\|^2 = \|\mathbf{x} - \mathbf{c}_k - \alpha \mathbf{a}_k\|^2$$

where

$$\alpha = (\mathbf{x} - \mathbf{c}_k)^T \mathbf{a}_k.$$

Here, a_k represents the first principal component vector found using PCA.

Plane Distance:

$$Dist(\mathbf{x}, C_k) = \|\mathbf{x}^\perp\|^2 = \|\mathbf{x} - \mathbf{c}_k - \alpha \mathbf{a}_k - \beta \mathbf{b}_k\|^2$$

where

$$\alpha = (\mathbf{x} - \mathbf{c}_k)^T \mathbf{a}_k, \quad \beta = (\mathbf{x} - \mathbf{c}_k)^T \mathbf{b}_k.$$

Here, a_k again represents the first principal component vector found using PCA, and b_k is the second principal component.

Sphere Distance:

$$\min_{\mathbf{c}_k} \sum_{i \in C_k} \max \left(0, \|\mathbf{x}_i - \mathbf{c}_k\|^2 - \eta \sum_{j \in C_k} \|\mathbf{x}_j - \mathbf{c}_k\|^2 \right)$$

For our synthetic data, we wanted to demonstrate the ability to cluster for a variety of shapes. Data includes 1D lines, 2D planes, and 3D spheres, and the goal is for K-subspaces to cluster those shapes together separately even when they are close together.

4. Datasets Used
5. Performance